

Bekzat Uraimov

Bachelor of Science in Computer Science, Bellevue College · Backend with Python · Building systems in C++
Seattle, WA · bkzturaimov@gmail.com · [bekzat.dev](#) · [github.com/bekzat-uraimov](#) · [linkedin.com/in/bekzat-uraimov](#)

EDUCATION

Bellevue College, Bachelor of Science in Computer Science <i>Current courses: Data Structures (C++), Python for Data Science</i>	Expected Apr 2028
Cascadia College, Associate Degree, Direct Transfer (DTA) · GPA 3.6 · 3 Java courses	Completed Aug 2026

EXPERIENCE

ONER , Software Engineer (part-time) <i>Built the backend of an online filmmaking course platform for Central Asia, preparing for launch · Python, FastAPI, PostgreSQL, Docker · code</i>	May 2026 – Present · Remote
- Stopped double charges and free access by checking every payment webhook 3 ways (signature, amount, duplicate) before a course opens, and returning the same order when someone clicks "buy" twice. - Kept paid lessons locked by checking access on the server against one table of who owns what, so the browser can never unlock a course by itself. - Made paid videos hard to share by playing them through Kinescope DRM, with short-lived tokens only owners get. - Let the team run the platform without touching the database by building a CRUD admin API for courses, uploads, users, purchases and refunds.	
Ilm-failover , Research project (in progress) <i>Fault tolerant LLM inference · C++ workers on llama2.c, Python router, TCP · github.com/bekzat-uraimov/Ilm-failover</i>	Oct 2026 – Present
- Building a way to keep an AI answer going when the machine writing it crashes: another worker rebuilds the state from a saved token log and finishes the same answer, byte for byte. - Keeps the state that must survive a crash about 3,000x smaller than the model's KV cache (about 1 KB vs 3.5 MB) by saving only the tokens and the random number state, then replaying them on the new worker.	
Akyldoo AI , Software Engineer (part-time) <i>Student startup building ThinkCoder, an AI coding assistant · Python, LiteLLM, Ollama</i>	Mar 2026 – May 2026 · Seattle, WA
Cut estimated AI API spend by about 40% by sending easy requests to a free local model and only hard ones to a paid cloud model, using LiteLLM.	
Video Editor and Colorist <i>Metr Production Studio, Antares, Mostcreative Agency · commercial videos, from shoot to final color grade</i>	Mar 2019 – Feb 2023 · Kyrgyzstan

HACKATHONS

Poly Predictor Kit · Polymarket Track Winner , QuackHacks, University of Oregon	Nov 2025
Led a team of 6 to build a Chrome extension that sums up any Polymarket event and points out the risks.	
AI Visual Novel Creator · Best Use of AI , CodeDay Fall 2025, Seattle	Nov 2025
Led a team of 6 to turn one prompt into a playable visual novel.	

SKILLS

Languages	Python, C++, Java, SQL
Backend	FastAPI, Flask, PostgreSQL, REST APIs, JWT auth, pytest
AI	LiteLLM, Ollama
Tools	Docker, Git, GitHub Actions, Cloudflare R2